

Measuring Morality in Political Attitude Expression*

Patrick W. Kraft[†]

February 22, 2017

Abstract

This study explores whether and how individuals evoke moral considerations when discussing their political beliefs. Analyzing open-ended responses in the 2012 American National Election Study using a previously validated dictionary, I find systematic ideological differences in moral reasoning—even when respondents are not explicitly asked about morality. The study proceeds to show that the reliance on moral considerations in attitude expression is conditional on the moral content of individual media environments.

*An earlier version of this paper was presented at the 73rd Annual Conference of the Midwest Political Science Association, April 16-19, 2015. The manuscript and code are available on GitHub: <https://github.com/pwkraft/mft>.

[†]Ph.D. Candidate, Stony Brook University, patrick.kraft@stonybrook.edu.

The increasing levels of polarization in contemporary politics renewed scholarly interest in psychological and attitudinal differences between liberals and conservatives (Jost, 2006). One such area of research focuses on the moral underpinnings of ideology. According to *Moral Foundations Theory* (MFT), moral thinking is organized by five innate intuitions: harm/care, fairness/reciprocity, ingroup/loyalty, authority/respect, and purity/sanctity (Haidt, 2012; Graham et al., 2013). Liberals and conservatives differ in their relative emphasis on these foundations, with liberals prioritizing the foundations of harm/care and fairness/reciprocity, and conservatives endorsing all five dimensions equally (Graham, Haidt, and Nosek, 2009).

A series of studies builds on MFT and provides evidence for the influence of moral foundations on issue preferences (Koleva et al., 2012; Kertzer et al., 2014), candidate trait evaluations (Clifford, 2014), and vote choice (Iyer et al., 2010; Franks and Scherr, 2015). Recent research further suggests that framing, campaigns, and elite communication can affect the relevance of individual moral foundations for attitude formation and attitude change (e.g. Clifford and Jerit, 2013; Clifford, Jerit, Rainey, and Motyl, 2015; Day et al., 2014; Feinberg and Willer, 2013). These analyses mostly relied on the Moral Foundations Questionnaire (MFQ) to measure moral reasoning, which explicitly asks respondents to judge the importance of considerations related to the five foundations (e.g. Graham et al., 2011). However, explicitly inquiring about moral arguments makes it impossible to examine directly how the connection between moral values and politics manifests itself in individual reasoning. Some scholars have criticized the MFQ because it does not ask people to make moral judgments per se (e.g. Clifford, Iyengar, Cabeza, and Sinnott-Armstrong, 2015). Indeed, Graham, Haidt, and Nosek (2009, 1031) describe the reports on moral relevance as “self-theories of moral judgment,” rather than direct measures of judgment itself. Such abstract self-theories might, in turn, deviate from actual judgments in specific situations.

The present study addresses this gap by examining whether people utilize the foundations in their day-to-day political reasoning (i.e., without being prompted by the language of a questionnaire). Moving beyond the MFQ and measuring moral reasoning as the explicit

expression of foundations allows us to examine the degree to which political attitudes are influenced by morality in a more unobtrusive context. The measurement approach proposed here uses a moral dictionary validated in previous studies (e.g., [Graham, Haidt, and Nosek, 2009](#)) to analyze individual verbatim responses to open-ended questions about political attitudes and preferences. Using open-ended items in such a way allows for a direct investigation of moral reasoning where the potential connection between morality and politics is not induced or facilitated by design. Furthermore, analyzing how individuals evoke moral considerations when they talk about politics is essential to understand the role of morality in political discussions and social influence.

Insofar as moral intuitions play a role in political reasoning, citizens should rely on the moral foundations when reporting their attitudes towards political actors, even if they are not explicitly asked to do so. Thus, the first step of the analyses focuses on the replication of previous findings connected to MFT and ideology using open-ended survey responses in the 2012 ANES. Consistent with MFT, the results reveal systematic differences between liberals and conservatives in the reliance on specific moral considerations even without being cued to think about morality. Furthermore, these differences in moral reasoning are shown to be politically consequential since they influence candidate preferences and vote choice—even after controlling for a person’s party identification. Integrating a large-scale content analysis of individual media environments, the analyses proceed to show that individuals who are exposed to moral rhetoric in political news are more likely to rely on moral considerations when discussing their own political beliefs. From a methodological standpoint, the present study demonstrates how open-ended survey responses can be utilized to investigate the underpinnings of political reasoning. Specifically, I introduce methods to improve conventional dictionary-based approaches for the analysis of open-ended responses and showcase the integration of media content analyses to trace the influence of exposure to political discourse on individual response behavior.

Measuring Moral Reasoning in Open-Ended Responses

The present study is the first to investigate moral reasoning by examining individual verbatim expressions of political attitudes and preferences. Using the moral foundations dictionary created by [Graham, Haidt, and Nosek \(2009\)](#), I identify references to specific moral considerations when respondents discuss aspects that they like and dislike about political parties and candidates. While the dictionary was initially used to analyze the content of liberal and conservative sermons, it was developed as a collection of general associations related to each foundation that could be applied in any context.¹ Other studies in the domain of MFT relied on (variations of) the dictionary to investigate moral considerations in elite communication (e.g. in news media coverage about stem cell research, [Clifford and Jerit, 2013](#)), but to date no research has examined individual political reasoning in open-ended survey responses.

Based on the terms signaling each foundation in the dictionary, any document can be scored according to its emphasis on the respective moral dimension. Conventional dictionary-based methods usually consist of the proportion of signal word occurrences in each document (e.g. [Graham, Haidt, and Nosek, 2009](#)). However, some terms in the dictionary might be problematic when applied to verbatim survey responses. In particular, when respondents describe their attitudes towards political actors, certain words might be too ubiquitous to be regarded as an unambiguous indicator for specific moral considerations. For example, the moral foundations dictionary includes “leader” as a signal word for the authority/respect dimension. What if almost all respondents are inclined to describe the qualities of presidential candidates as *leaders*, irrespective of explicit moral considerations related to authority?

One way to address this problem using conventional methods would be to revise the dictionary and eliminate words that are deemed problematic. However, this leaves a lot of discretion to the researcher. Drawing on techniques developed in the field of information retrieval, I propose an alternative approach. If it is the case that “leader” represents a term that is commonly used to describe presidential candidates (irrespective of moral considera-

¹See Appendix ?? for the full content of the dictionary.

tions), then it should appear more frequently in open-ended responses across all individuals. Terms that are used by almost all respondents therefore provide less information about differences in their (moral) reasoning than terms that only occur in few responses. Stated differently, if a specific moral word is mentioned by a large majority of respondents, it is more likely that the term can be used in multiple contexts and is therefore not necessarily unique to the moral domain. In the analyses below, MFT scores are computed for a foundation by weighting each term in the dictionary according to its ubiquity across documents, which serves as a proxy for the term’s discriminative information:

$$\text{MFT}_{if} = \frac{1}{W_i} \sum_{t \in \mathcal{D}_f} \left[w_{it} * \log_{10} \left(\frac{N}{n_t + 1} \right) \right], \quad (1)$$

where MFT_{if} denotes the score of document i for foundation f , W_i is the total number of words in document i , t indicates a term in the set of signal terms in foundation dictionary \mathcal{D}_f , and w_{it} denotes the number of occurrences of term t in document i . Furthermore, N denotes the total number of documents, and n_t is the number of documents in which the term t appears. The weight represents the inverse of the proportion of documents in which the target term appears.² As such, terms that are ubiquitous across the entire corpus receive a lower weight, and terms that appear in only few documents receive a higher weight. The denominator in equation (1) includes +1 to ensure that it does not equal zero if a dictionary term does not appear in any of the documents.

In the analyses presented here, each document is an individual’s verbatim response to a set of open-ended questions. As such, a respondent’s MFT score for foundation f is the weighted proportion of words in the response that signal the respective foundation. The score has a lower bound of 0 (document does not contain any dictionary terms) and is independent of document length (since it is based on relative occurrences). Higher scores

²This specification is usually referred to as tf-idf weighting and is commonly used in quantitative text analyses. The acronym tf-idf stands for “term frequency - inverse document frequency,” which refers to the rationale that the frequency of specific terms are weighted by the inverse of the frequency of occurrence across documents. See [Manning et al. \(2008, ch. 6\)](#) for an introduction.

imply larger proportions of dictionary terms in a document. Most importantly, however, words that appear in nearly all open-ended remarks affect MFT scores less than the words mentioned only by few respondents because ubiquitous words convey less information about differences across individuals. Overall, the MFT score provides a correction for potential distortions due to suboptimal terms in the dictionary while leaving its exact content outside of the researcher’s discretion. Since nominal values of the MFT score above zero do not have a clear substantive interpretation, they are rescaled to unit variance.

Results

The analyses are based on the 2012 ANES. The primary dependent variables (i.e., the MFT scores described above) are based on open-ended questions in which respondents were asked to report what they *liked* and *disliked* about either presidential candidate as well as the Republican and Democratic parties. The responses were aggregated for each individual and pre-processed by correcting spelling errors using an implementation of the Aspell spell checking algorithm in R (www.aspell.net).³ Due to the fact that the MFT scores are bound at zero (i.e., none of the words in the dictionary appear in the response), individual response patterns are modeled via Tobit regressions for each of the moral foundations under consideration. The Tobit framework allows us to decompose the estimates into the effect on the probability of mentioning a specific foundation *at all* as well as the degree of emphasis on the foundation given that it was mentioned by a respondent (see McDonald and Moffitt, 1980, for details on the decomposition of Tobit model estimates).

³Please refer to the appendix for further information on the data and recoding. Note that the dimension of purity/sanctity was omitted from the analyses of open-ended responses due to its low general prevalence in individual attitude expressions.

Ideological Differences in Moral Reasoning

We begin by estimating a set of Tobit regressions using ideology (and the control variables discussed above) to predict the individual MFT score for each of the moral foundations (excluding purity/sanctity).⁴ To reiterate, the MFT score measures the weighted proportion of moral foundation terms in an open-ended response. Figure 1 compares liberals and conservatives while holding all other variables constant at their respective means. The effects of ideology (liberal - conservative) estimated in the Tobit model are decomposed into two parts: the left panel displays the change in probability of mentioning a specific foundation at all (i.e., probability of the MFT score to be larger than zero), whereas the right panel displays the expected change in the degree of emphasis on a foundation given that it was mentioned (i.e., the change in the MFT score given that it is larger than zero, measured in standard deviations).




Figure 1: Difference between liberals and conservatives in the probability of mentioning each moral foundation (left panel), and in the MFT score given that the foundation was mentioned (right panel), holding all other control variables at their respective means (along with 95% confidence intervals). Positive values indicate that liberals are more likely to mention the respective foundation than conservatives (left panel), or emphasize it more than conservatives (right panel), and vice versa. Estimates are based on separate Tobit models for each foundation's MFT score. The dimension of purity/sanctity was omitted due to its low general prevalence in individual attitude expressions. Control variables include church attendance, education, age, sex, race, survey mode, response length, and the Wordsum vocabulary score. Full model results are displayed in the appendix, Table ??.

⁴The full estimates for this and all subsequent models discussed the remainder of the paper are presented in Appendix ??.

Positive values denote a higher probability of mentioning the respective moral foundation (left panel) or a higher MFT score (right panel) of a response among individuals who identified as liberals, while negative values indicate a higher probability/higher score among conservatives. The effects are consistent with the the expectations of MFT for three out of four moral foundations. Liberals are significantly more likely to mention the foundations of harm/care and fairness/reciprocity. For respondents who identified as liberal as compared to conservative, the probability of referencing these two foundations is increased by about 6 percentage points. Furthermore, given that respondents mention these two foundations at all, liberals emphasize it more than conservatives when evaluating political parties and candidates. The MFT score for the harm/care foundation is about 0.07 standard deviations higher among liberals than conservatives. The effect is slightly larger for the fairness/reciprocity dimension. Conversely, being conservative increases the MFT score for the foundation of ingroup/loyalty by about 0.09 standard deviations. For the authority/respect dimension, we do not observe any significant differences between liberals and conservatives.

The Political Relevance of Moral Reasoning

A skeptical reader may argue that even if the ideological patterns are consistent with MFT, the expression of moral considerations might not be as strongly related to other forms of political behavior as latent moral foundations measured by the MFQ. Previous research relying on the MFQ has linked moral foundations to an array of political outcomes, such as candidate preferences (Iyer et al., 2010) and voting behavior (Franks and Scherr, 2015). Figure 2 presents the changes in expected probabilities of voting for the Democratic (vs. the Republican) presidential candidate in the 2012 election for individuals emphasizing the moral foundations in their open-ended responses. The estimated probabilities are based on logit models including MFT scores for each moral foundation as independent variables (as well as the remaining controls), which were held constant at their mean values when computing expected values. Individuals who emphasized moral considerations related to the harm/care

and fairness/reciprocity foundations are more likely to vote for Barack Obama than for Mitt Romney. Respondents who emphasized the ingroup/loyalty foundation, on the other hand, were less likely to vote for Obama.⁵




Figure 2: Change in predicted probabilities to vote for the Democratic rather than Republican candidate when MFT score is increased from its minimum (no overlap between dictionary and response) by one standard deviation, holding all other control variables constant at their respective means (along with 95% confidence intervals). Positive values indicate that respondents who emphasize the respective foundation are more likely to vote for the Democratic candidate, and vice versa. Estimates are based on a single logit model including MFT scores or each foundation. The dimension of purity/sanctity was omitted due to its low general prevalence in individual attitude expressions. Control variables include party identification, church attendance, education, age, sex, race, survey mode, response length, and the Wordsum vocabulary score. Full model results are displayed in the appendix, Table ??.

The effects on vote choice might not seem large, but bear in mind that the measure of moral reasoning is based solely on the content of open-ended responses in which respondents were *not* explicitly asked about morality. Yet, the moral considerations evoked by respondents were powerfully related to party and candidate evaluations as well as vote choice, even after controlling for strength of party identification, the length of individual responses, and general verbal skills. Overall, the analyses show that people’s open-ended comments about both candidates and both parties are imbued with moral content and that these comments relate to political judgments in the manner predicted by MFT.

⁵The same results hold when looking at the effect on feeling thermometers towards both parties and candidates. Please refer to the appendix for details.

Media Content and Exposure to Moral Rhetoric

Next, we investigate whether the general reliance on moral considerations is a product of exposure to moralized political discourse. For each individual, I compute the sum of MFT scores to measure the general tendency to emphasize *any* moral foundation. The main independent captures moralization of media environments based on a content analyses of media sources used by each individual. The 2012 ANES includes a large array of items indicating whether respondents regularly consumed various news outlets. For all media sources available, I downloaded the content of the coverage on either presidential candidates during the survey field period in the last month of the campaign (October 2012) from Lexis-Nexis and coded their emphasis on moral foundations using the same approach as for open-ended survey responses. Similar to the respondent MFT scores, the media scores were summed over all foundations to capture the general degree of moralization.⁶

Based on the coded content for each media source, I create a measure that represents the extent to which each individual’s media environment emphasized moral considerations. For each respondent in the ANES, I select the media sources he or she reported to watch/read regularly and computed the sum of the sources’ MFT scores. Using this approach, I can analyze whether individuals who rely on media sources that use more moralized reporting were also more likely to emphasize moral arguments in their open-ended responses.

The results are presented in Figure 3. Estimates are based on Tobit models that take into account the censoring of the moral reasoning measure and effects are decomposed into the probability of mentioning any moral foundation (left panel) as well as the emphasis on morality, given that any foundation was mentioned (right panel). Individuals who are exposed to media sources that report on the campaign in a more moralized manner put a stronger emphasis on moral arguments in their open-ended responses describing their political attitudes about the parties and candidates. Thus, citizens learn to embed moral reasoning in

⁶In total, I retrieved the content of 28 media sources, such as the New York Times (print and online), CNN.com, or various Fox News Programs. See Figure ?? in the Appendix for a more detailed overview of the media outlets and their respective MFT scores.




Figure 3: Effect of MFT content in individual media environments on the probability of mentioning any moral foundation (left panel), and on the summed MFT score given that any foundation was mentioned (right panel), holding all other control variables at their respective means (along with 95% confidence intervals). For each moral dimension, the independent variable consists of the sum of MFT scores for media sources that are regularly watched/read by each respondent. Positive values indicate that individuals who consume media content that on average emphasizes the respective moral foundation more than other sources, are more likely to mention the respective foundation (left panel), or emphasize it more (right panel), and vice versa. Estimates are based on a Tobit model. Control variables include general media exposure, political knowledge, political discussion frequency, church attendance, education, age, sex, race, survey mode, response length, and the Wordsum vocabulary score. Full model results are displayed in the appendix, Table ??.

their political evaluations by adopting moral arguments from their media environment. It is worth emphasizing again that the model presented here control for education, logged overall response length, and the Wordsum score, which should account for potential confounding factors related to the respondents’ eloquence when discussing their political attitudes.

Robustness Checks

The results discussed thus far rely on the crucial assumption that the dictionary-based approach for open-ended responses captures the theoretical concept of interest—*moral* reasoning. The remaining section briefly considers the possibility that this underlying assumption was not met and that the ideological divide observed throughout the analyses does not reflect differences in expressed moral foundations. Indeed, the terms in the dictionary may coincidentally recover unrelated differences in word choice between liberals and conservatives

when discussing their attitudes towards parties and candidates in the 2012 U.S. Presidential election. For example, one prominent issue in the election was the Affordable Care Act, which might increase the likelihood of Democrats mentioning the term “care” and thereby increasing the emphasis on the harm/care foundation irrespective of underlying moral considerations. As such, the observed differences might be an artifact due to the nature of the questions under considerations as well as the specific political context of the presidential campaign.

In order to address these concerns, I replicated the first model using open-ended responses from a survey administered in a different political context. The survey was conducted via telephone with 594 adults aged 18 or older between early January, 2001 and July, 2003. The telephone numbers were a random-digit-dial (RDD) sample drawn from residents within a 25 mile radius of a large northeastern state university. As such, the survey was not conducted during a major presidential election campaign. Furthermore, the survey uses a different set of open-ended items. Rather than asking about attitudes towards presidential candidates and both major parties, respondents were asked to describe liberals and conservatives as well as their respective beliefs in general. The coding and analyses are equivalent to those for Figure 1, although the survey did not contain the Wordsum scores (and varying survey mode) included as controls in the main analyses. The results are displayed in Figure 4.

The patterns are strikingly consistent with previous results. Liberals are more likely to emphasize the foundations of harm/care and fairness/reciprocity. The result for the ingroup/loyalty dimension, however, do not reach conventional levels of statistical significance. Additional analyses reveal that the ideological differences in moral reasoning are mostly due to the fact that individuals who identify as liberals emphasize the foundations of harm/care and fairness/reciprocity more strongly than conservatives when describing their ingroup (i.e., other liberals and their beliefs), while conservatives emphasize the foundation of ingroup/loyalty more strongly than liberals when describing their ingroup (results available upon request). The fact that the same basic ideological pattern can be recovered in




Figure 4: Replication of main model (c.f., Figure 1) using RDD sample from residents within a 25 mile radius of a large northeastern state university. Figure displays difference between liberals and conservatives in the probability of mentioning each moral foundation (left panel), and in the MFT score given that the foundation was mentioned (right panel), holding all other control variables at their respective means (along with 95% confidence intervals). Positive values indicate that liberals are more likely to mention the respective foundation than conservatives (left panel), or emphasize it more than conservatives (right panel), and vice versa. Estimates are based on separate Tobit models for each foundation’s MFT score. The dimension of purity/sanctity was omitted due to its low general prevalence in individual attitude expressions. Control variables include church attendance, education, age, sex, race, and response length. Full model results are displayed in the appendix, Table ??

a survey that was conducted in a different political context (non-election period, Republican administration), employed a different survey mode (phone interview), and relied on a different set of open-ended survey questions (asking about liberals and conservatives and their respective beliefs), provides further support for the argument that the MFT dictionary allows us to recover basic moral considerations in political reasoning.

A related concern might be the question whether the content analysis of media sources using the dictionary is able to capture overall levels of moralization in news reporting. Luckily, a study reported in [Feinberg and Willer \(2013\)](#) included manual coding a selection of newspaper articles on environmental issues to capture whether they use rhetoric grounded in each of the moral domains. Their coding therefore focuses on the same foundations without utilizing the dictionary. I computed a general moralization variable by summing the scores used in [Feinberg and Willer \(2013\)](#) and compared them to the MFT scores based on the

procedures outlined above.⁷




Figure 5: Validity check based on the data from [Feinberg and Willer \(2013\)](#).

Figure 5 presents the correlation of general moralization in each article based on the manual coding in [Feinberg and Willer \(2013\)](#) compared to the dictionary method used in the analyses presented here. While the correlation is far from being perfect, the weighted dictionary method clearly captures some of the same variance as manual assessments of the emphasis on moral foundations.

Discussion

There is much more to be learned about the relationship between moral foundations and broader political attitudes and ideology. To fully understand how political views are rooted in morality, researchers need to take a closer look at the conditions under which citizens express moral considerations in the context of politics. This study utilized open-ended survey responses to investigate morality in individual attitude expression. The analyses of open-ended survey responses is valuable in this context because it allows us to evaluate whether citizens make references to moral considerations in a political context that does not induce an explicit connection to morality. As such, it can be directly investigated when and how ideological differences in the emphasis of moral foundations manifest themselves in individual reasoning about political actors. More generally, open-ended survey responses

⁷I am indebted to the authors for providing the data.

provide a promising and still largely neglected data source to investigate the determinants and structure of ideology and political reasoning. In particular, scholars can directly assess moral reasoning in surveys that do not contain the MFQ or related measures, simply by relying on open-ended survey responses. More broadly, focusing on open-ended measures provides new opportunities to study the role of morality in day-to-day political reasoning, discussions, and social influence.

The empirical analyses presented here extend and qualify previous research on moral foundations and ideology. The results showed systematic patterns in the emphasis on moral considerations among liberals and conservatives consistent with MFT for three out of four foundations. Liberals are more likely to mention considerations related to harm/care and fairness/reciprocity when discussing their political preferences, whereas conservatives are more likely to emphasize the moral foundation of ingroup/loyalty. The second part of the analyses focused on the political relevance of moral reasoning as conceptualized by open-ended survey responses. Here, the results revealed consistent relationships between individual moral foundations and voting behavior, which showed that moral reasoning (measured via open-ended survey responses) is a politically meaningful and influential concept. Lastly, exposure to moralized political discourse increase the reliance on moral considerations.

The present study therefore reaffirms the importance of moral reasoning in politics but also reveals its potential conditionality. Ultimately, a deeper understanding of the role of morality in politics necessitates further analyses of the broader political context that shapes individual information environments (e.g., how the endorsement of moral foundations varies over time and across campaigns). Such an investigation would further illuminate how exposure to political discourse fosters ideological differences in moral reasoning. In times of growing partisan polarization, a better understanding of the antecedents of this ideological divide is essential.

References

- Clifford, Scott. 2014. "Linking Issue Stances and Trait Inferences: A Theory of Moral Exemplification." *Journal of Politics* 76 (3): 698–710.
- Clifford, Scott, and Jennifer Jerit. 2013. "How Words Do the Work of Politics: Moral Foundations Theory and the Debate over Stem Cell Research." *Journal of Politics* 75 (3): 659–671.
- Clifford, Scott, Jennifer Jerit, Carlisle Rainey, and Matt Motyl. 2015. "Moral Concerns and Policy Attitudes: Investigating the Influence of Elite Rhetoric." *Political Communication* 32 (2): 229–248.
- Clifford, Scott, Vijeth Iyengar, Roberto Cabeza, and Walter Sinnott-Armstrong. 2015. "Moral Foundations Vignettes: A Standardized Stimulus Database of Scenarios Based on Moral Foundations Theory." *Behavior Research Methods* 47 (4): 1178–1198.
- Day, Martin V., Susan T. Fiske, Emily L. Downing, and Thomas E. Trail. 2014. "Shifting Liberal and Conservative Attitudes using Moral Foundations Theory." *Personality and Social Psychology Bulletin* 40 (12): 1559–1573.
- Feinberg, Matthew, and Robb Willer. 2013. "The moral roots of environmental attitudes." *Psychological Science* 24 (1): 56–62.
- Franks, Andrew S., and Kyle C. Scherr. 2015. "Using Moral Foundations to Predict Voting Behavior: Regression Models from the 2012 US Presidential Election." *Analyses of Social Issues and Public Policy* 15 (1): 213–231.
- Graham, Jesse, Brian A. Nosek, Jonathan Haidt, Ravi Iyer, Spassena Koleva, and Peter H. Ditto. 2011. "Mapping the Moral Domain." *Journal of Personality and Social Psychology* 101 (2): 366–385.

- Graham, Jesse, Jonathan Haidt, and Brian A. Nosek. 2009. "Liberals and Conservatives Rely on Different Sets of Moral Foundations." *Journal of Personality and Social Psychology* 96 (5): 1029–1046.
- Graham, Jesse, Jonathan Haidt, Sena Koleva, Matt Motyl, Ravi Iyer, S. Wojcik, and Peter H. Ditto. 2013. "Moral Foundations Theory: The pragmatic Validity of Moral Pluralism." *Advances in Experimental Social Psychology* 47: 55–130.
- Haidt, Jonathan. 2012. *The Righteous Mind: Why Good People are Divided by Politics and Religion*. New York: Random House.
- Iyer, Ravi, Jesse Graham, Spassena Koleva, Peter Ditto, and Jonathan Haidt. 2010. "Beyond Identity Politics: Moral Psychology and the 2008 Democratic Primary." *Analyses of Social Issues and Public Policy* 10 (1): 293–306.
- Jost, John T. 2006. "The End of the End of Ideology." *American Psychologist* 61 (7): 651–670.
- Kertzer, Joshua D., Kathleen E. Powers, Brian C. Rathbun, and Ravi Iyer. 2014. "Moral Support: How Moral Values Shape Foreign Policy Attitudes." *Journal of Politics* 76 (3): 825–840.
- Koleva, Spassena P., Jesse Graham, Ravi Iyer, Peter H. Ditto, and Jonathan Haidt. 2012. "Tracing the Threads: How Five Moral Concerns (Especially Purity) Help Explain Culture War Attitudes." *Journal of Research in Personality* 46 (2): 184–194.
- Manning, Christopher D., Prabhakar Raghavan, Hinrich Schütze et al. 2008. *Introduction to Information Retrieval*. Cambridge: Cambridge University Press.
- McDonald, John F., and Robert A. Moffitt. 1980. "The Uses of Tobit Analysis." *The Review of Economics and Statistics* 62 (2): 318–321.